

# Introduction to Deep Reinforcement Learning

[https://github.com/racousin/rl\\_introduction](https://github.com/racousin/rl_introduction)

Raphael Cousin

January 22, 2023

# Course Objective

- The keys to go by yourself in RL
- Practice coding
- General culture

What do you already know about RL?

## 0.1)Intro



Figure 1: 2017, AlphaGo beat Ke Jie, the number one ranked player in the world


## 0.2) RL!=ML?

- Supervised learning:  $(X, Y) \rightarrow \hat{f}(X) = Y$
- Unsupervised learning:  $X \rightarrow \hat{f}(X|\hat{Y})$
- Reinforcement learning (RL): environment/goal  $\rightarrow$  optimally interact with the environment to achieve the goal

## 0.3) Agent takes action in an Environment

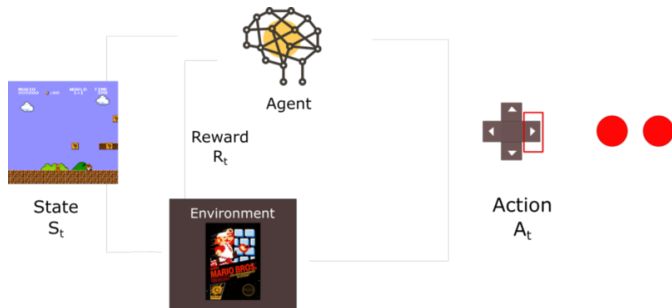
RL framework:

- Unknown “correct” actions → no explicit supervision for a learning algorithm to try to mimic.
- Provide our algorithms only a reward function → indicates to the learning agent when it is doing well, and when it is doing poorly.



## 0.4) Application

Applications: Autonomous vehicle, robot legged locomotion, cell-phone network routing, marketing strategy, factory control, efficient web-page indexing, video games, Neural Architecture Search, ...



- I) MDP - The RL framework
- II) RL - Model free
- III) Deep RL - Model free

# I) MDP and dynamic Programming

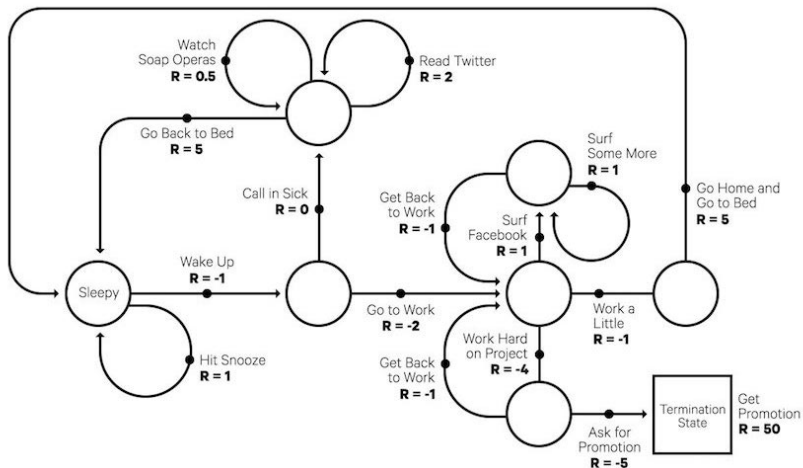


# I.1) Markov Decision Process

- State space:  $S$
- Action space:  $A$
- Transition model:  $P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$
- Immediate reward  $R_a(s, s') = R_{t+1} \in \mathbb{R}$
- Policy function  $\pi(s) \in A$

$\forall s, s \in S, a \in A, t \in \mathbb{N}$

## I.2) Example: MDP long term reward



## I.3) The MDP Objective

- Find the optimal policy:

$$\pi^* = \arg \max_{\pi} E[G]$$

- Finite-horizon return:  $G_t = \sum_{k=0}^T R_{t+k+1}$ .

- Infinite-horizon discounted return:  $G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$

$$\gamma \in ]0, 1[, t \in \mathbb{N}$$

## I.4) Trajectories

- Trajectories:  $\tau = (s_0, a_0, s_1, a_1, \dots)$
- States follow:  
 $s_0 \sim \rho_0(\cdot)$   
 $s_{t+1} \sim P(\cdot | s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = P(\cdot | s_t, a_t)$
- Actions follow:  $s_t \sim \pi(s_t)$
- So probability of a T -step trajectory is:  
 $\mathbb{P}(\tau | \pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1} | s_t, a_t) \pi(a_t | s_t)$
- $\pi^* = \arg \max_{\pi} E_{\tau \sim \pi} [G(\tau)] = \int_{\tau} \mathbb{P}(\tau | \pi) G(\tau)$

# Coding session

## Environment and Agent

## I.5) Value and Action-Value Functions

- Value Function:  $V^\pi(s) = E_{\tau \sim \pi}[G_t | s_t = s]$
- Action-Value Function:  
 $Q^\pi(s, a) = E_{\tau \sim \pi}[G_t | s_t = s, a_t = a]$
- The Optimal Value Function:  
 $V^*(s) = \max_\pi E_{\tau \sim \pi}[G_t | s_t = s] = E_{\pi^*}[G_t | s_t = s]$
- Optimal Action-Value Function:  
 $Q^*(s, a) = \max_\pi E_{\tau \sim \pi}[G_t | s_t = s, a_t = a] = E_{\pi^*}[G_t | s_t = s, a_t = a]$

## I.6) Bellman Equations

- **Idea :** The value of your starting point is the reward you expect to get from being there, plus the value of wherever you land next.

$$V(s) = \mathbb{E}[G_t | S_t = s]$$

$$= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

$$Q(s, a) = \mathbb{E}[R_{t+1} + \gamma \mathbb{E}_{a \sim \pi} Q(S_{t+1}, a) \mid S_t = s, A_t = a]$$

## I.7) Bellman Equations development

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) Q_{\pi}(s, a)$$

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s')$$

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left( R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s') \right)$$

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') Q_{\pi}(s', a')$$



## I.8) Bellman Equations Optimality

Bellman equations for the optimal value functions

$$V_*(s) = \max_{a \in \mathcal{A}} Q_*(s, a)$$

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_*(s')$$

$$V_*(s) = \max_{a \in \mathcal{A}} \left( R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_*(s') \right)$$

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \max_{a' \in \mathcal{A}} Q_*(s', a')$$

## I.9) The MDP Solution

Dynamic Programming allows to resolve the MDP optimization problem ( $\pi^* = \arg \max_{\pi} E_{\tau \sim \pi}[G(\tau)]$ ). It is an iterative process:

- Policy initialization
- Policy evaluation
- Policy improvement

## I.10) Policy evaluation

Policy Evaluation: compute the state-value  $V_\pi$  for a given policy  $\pi$ :  
We initialize  $V_0$  arbitrarily. And we update it using:

$$\begin{aligned} V_{k+1}(s) &= \mathbb{E}_\pi[r + \gamma V_k(s_{t+1}) | S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} P(s', r|s, a) (r + \gamma V_\pi(s')) \quad (1) \end{aligned}$$

$V_\pi(s)$  is a fix point for (1), so if  $(V_k)_{k \in \mathbb{N}}$  converges, it converges to  $V_\pi$ .

## I.11) Policy Improvement

Policy Improvement: generates a better policy  $\pi' \geq \pi$  by acting greedily. Compute  $Q$  from  $V$  ( $\forall a, s$ ):

$$\begin{aligned} Q_{\pi}(s, a) &= \mathbb{E}[R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t = s, A_t = a] \\ &= \sum_{s', r} P(s', r | s, a) (r + \gamma V_{\pi}(s')) \end{aligned}$$

Update greedily:  $\pi'(s) = \arg \max_{a \in \mathcal{A}} Q_{\pi}(s, a)$  ( $\forall s$ )

## I.12) Dynamic Programming

Policy Iteration: iterative procedure to improve the policy when combining policy evaluation and improvement.

$$\pi_0 \xrightarrow{\text{evaluation}} V_{\pi_0} \xrightarrow{\text{improve}} \pi_1 \xrightarrow{\text{evaluation}} \dots \xrightarrow{\text{improve}} \pi_* \xrightarrow{\text{evaluation}} V_* \quad (1)$$

# Take home message

Initialize  $\pi(s), \forall s$

- 1 Evaluate  $V_\pi(s), \forall s$  (using  $\mathbb{P}_{ss'}^a$ )
- 2 Compute  $Q_\pi(s, a), \forall s, a$  (using  $\mathbb{P}_{ss'}^a$ )
- 3 Update  $\pi'(s) = \max_a Q_\pi(s, a), \forall s$
- 4 While  $\pi'(s) \neq \pi(s)$  do  $\pi(s) = \pi'(s)$  and iterate

Result :  $\pi = \arg \max_\pi E[G]$

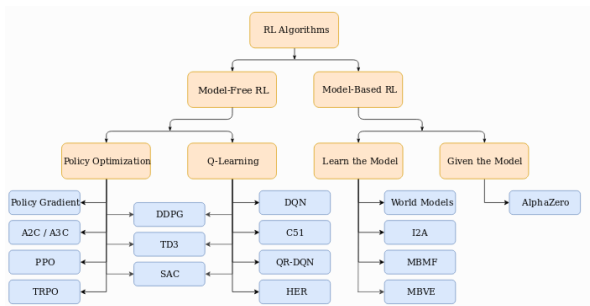
# Coding session

## Dynamic Programming





# I.18) Bestiary of RL Algorithms



To simplify:

- **Model free:** large data + low complexity
- **Model based:** less data + high complexity

## II.1) Model-Free

The objective is the same as MDP:

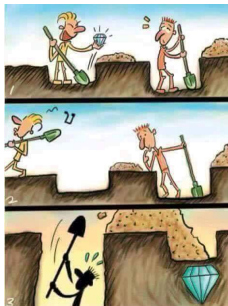
- $J(\pi) = \int_{\tau} \mathbb{P}(\tau|\pi)G(\tau) = E_{\tau \sim \pi}[G(\tau)]$
- The optimization problem:  
$$\pi^* = \arg \max_{\pi} J(\pi)$$

But we don't know the transition model  $P$ . The constraint is therefore to interact intelligently with the environment to obtain the information needed to solve the problem.

- **Q-learning:** learn the action value function  $Q$   
 $(\pi'(s) = \arg \max_{a \in \mathcal{A}} Q_{\pi}(s, a))$
- **Policy Optimization:** learn directly the policy  $\pi$

## II.2) Exploration-Exploitation

Knowledge of the environment comes from interaction. There are trade-offs to be made between using what we know and further exploration.



Naive solution: force exploration with random action (control by an  $\epsilon$  factor)

## II.3) Definition

- **On-policy:** Use the deterministic outcomes or samples from the target policy to train the algorithm.
- **Off-policy:** Training on a distribution of transitions or episodes produced by a different behavior policy rather than that produced by the target policy.

## II.4) Monte-Carlo

- To evaluate  $V_\pi(s) = E_{\tau \sim \pi}[G_t | s_t = s]$
- Generate an episode with the policy  $\pi$   $S_1, A_1, R_2, \dots, S_T$  to compute  $G_t = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1}$ .
- The empirical value function is :  $V_\pi(s) = \frac{\sum_{t=1}^T \mathbb{I}[S_t=s] G_t}{\sum_{t=1}^T \mathbb{I}[S_t=s]}$
- As, well, the empirical action-value function is :
- $Q_\pi(s, a) = \frac{\sum_{t=1}^T \mathbb{I}[S_t=s, A_t=a] G_t}{\sum_{t=1}^T \mathbb{I}[S_t=s, A_t=a]}$

## II.5) Monte-Carlo Algorithm

Initialize  $Q$   $Q(s, a) \forall s, a$ .

① Generate an episode with the policy  $\pi$  (extract from  $Q$   $\epsilon$ -greedy)

② Evaluate  $Q$  using the episode:

$$q_{\pi}(s, a) = \frac{\sum_{t=1}^T (\mathbb{I}[S_t=s, A_t=a] \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1})}{\sum_{t=1}^T \mathbb{I}[S_t=s, A_t=a]}$$

③ Improve the policy greedily:  $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$

④ Iterate

## II.6) Temporal difference/bootstrapping

- Remember  $V(S_t) = \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$
- So  $R_{t+1} + \gamma V(S_{t+1})$  is an unbiased estimate for  $V(S_t)$
- As well  $R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$  is an unbiased estimate for  $Q(S_t, A_t)$
- $R_{t+1} + \gamma V(S_{t+1})$  is called the TD target.
- $\alpha$  improvement:  
$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

This observation motivates the following algorithm.

## II.7) SARSA Algorithm

Initialize  $Q$  function  $Q(s, a) \forall s, a$

$S_t$  = initial state, act with  $\pi$  to get  $A_t, R_{t+1}, S_{t+1}$

- ① Act with  $\pi$  to get  $A_{t+1}, R_{t+2}, S_{t+2}$
- ② Evaluate  $Q$  using the observation step:  
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$
- ③ Improve the policy greedily:  $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$
- ④  $A_t = A_{t+1}, R_{t+1} = R_{t+2}, S_{t+1} = S_{t+2}$ . Iterate



## II.8) Q-learning

- Remember

$$Q_{\pi^*}(S_t, A_t) = \mathbb{E}[R_{t+1} + \gamma \max_{a'} Q_*(S_{t+1}, a') | S_t = s, A_t = a]$$

- So  $R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a')$  is an unbiased estimate for  $Q(S_t, A_t)$
- $R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a')$  is called the Q target.
- $\alpha$  improvement:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

This observation motivates the following algorithm.

## II.9) Q-learning Algorithm

Initialize  $Q$  function  $Q(s, a) \forall s, a$

$S_t = \text{initial state}$

- ①  $Q : S \times A \rightarrow \mathbb{R}$
- ②  $\forall s, a Q(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$
- ③ Improve the policy greedily:  $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$
- ④ Iterate

# Coding session

## Temporal Difference

## II.10) Policy Optimization

- Parametrization of policy,  $\pi_\theta$ .
- We aim to maximize the expected return  $J(\pi_\theta) = E_{\tau \sim \pi_\theta}[G(\tau)]$ .
- Gradient ascent:  
$$\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_\theta)|_{\theta_k}.$$
- We can proof that:  
$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\theta}[\sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) G(\tau)]$$
  
<https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>



## I.4) Introduction - Q Learning limitations

- $S$  high dimension  $\rightarrow$  low convergence
- It doesn't work for continuous action space

## II.1) Deep Q Learning

Parametrize  $Q$  with  $\theta$ , initialize  $\theta \in \mathbb{R}^d$

- $Q_\theta : S \times A \rightarrow \mathbb{R}$
- Objective find  $\theta^* \in \mathbb{R}^d \forall s, a \ Q_{\theta^*}(s, a) = \mathbb{E}_\pi^*[G_t | S_t = s, A_t = a]$
- $Q_{\pi^*}(S_t, A_t) = \mathbb{E}[R_{t+1} + \gamma \max_{a'} Q_*(S_{t+1}, a') | S_t = s, A_t = a]$
- $y = R_{t+1} + \gamma \max_{a'} Q_\theta(S_{t+1}, a')$  is called the Q target.
- Loss (eg MSE):  $L(\theta) = \mathbb{E}_{s, a \sim Q}[(y - Q(s, a, \theta))^2]$

## II.2) Deep Q Learning Algorithm

Initialize  $\theta \in \mathbb{R}^d$ , get initial state  $S_t$

- ①  $\forall a$  predict/compute  $Q_\theta(S_t, a)$
- ② Act greedily  $A_t = \arg \max([Q_\theta(s, a_0), Q_\theta(s, a_1), \dots, Q_\theta(s, a_{\dim(A)})])$   
and get  $R_{t+1}, S_{t+1}$
- ③  $\forall a$  predicts  $Q_\theta(S_{t+1}, a)$   
Compute the target  $y = R_{t+1} + \gamma \max_{a \in A} Q_\theta(S_{t+1}, a)$
- ④ Evaluate + improve  $Q$  by minimizing the loss  
 $L(\theta) = (y - Q(S_t, A_t, \theta))^2$
- ⑤ Iterate



## II.3) Convergence and stability improvements

### Experience replay

- At every step  $t$ , we get  $S_t, A_t, R_{t+1}, S_{t+1} \rightarrow$  memory
- New loss at each step

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[ \left( r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta_i) \right)^2 \right]$$

- Other improvements: Epsilon decay, Clipping, Double Q learning

# Coding session

## Deep Q - learning

## III.1) Policy Optimization

- Parametrize the policy,  $\pi_\theta$ .
- We aim to maximize the expected return  $J(\pi_\theta) = E_{\tau \sim \pi_\theta}[G(\tau)]$ .
- Gradient ascent:  
$$\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_\theta)|_{\theta_k}.$$
- We can proof that:  
$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\theta}[\sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) G(\tau)]$$
  
<https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>

## III.2) Reinforce/VPG algorithm

Initialize policy  $\pi_\theta$

- 1 Generate episodes  $\mathcal{D} = \{\tau_i\}_{i=1,\dots,N}$  with the policy  $\pi_\theta$
- 2 Compute gradient approximation
$$\hat{\nabla} = \frac{1}{|\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) G_t$$
- 3 Update policy (apply gradient ascent)  $\theta \leftarrow \theta + \alpha \hat{\nabla}$
- 4 Iterate

# Coding session

## Policy gradient - reinforce

### III.3) Improvement Actor-Critic

We can rewrite the policy gradient

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} [\sum_{t=0}^T \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \Phi_t],$$

whith  $\Phi_t$  could be any of

- $\Phi_t = G_t$
- $\Phi_t = \sum_{t'=t}^T R_{t+1} - V(s_t)$
- $\Phi_t = \sum_{t'=t}^T R_{t+1} - Q(s_t, a_t)$

For the last 2 cases we need to estimate V or Q (the critics).

## III.4) Actor-Critic Algorithm

Initialize  $\theta, \in \mathbb{R}^{d_1}, \phi \in \mathbb{R}^{d_2}$  Get start State  $S_t$

- ① Generate  $A_t$   $R_{t+1}$  following  $\pi_{\theta_t}$
- ② Update actor (apply gradient ascent)  
 $\theta \leftarrow \theta + \alpha \hat{\nabla} = \theta + Q_{\phi}(a_t, s_t) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$
- ③ Compute critic target  $y = R_{t+1} + \gamma Q_{\phi}(S_{t+1}, a')$
- ④ Evaluate/Improve  $Q_{\phi}$  by minimizing the loss  
 $L(\phi) = (y - Q(S_t, A_t, \phi))^2$
- ⑤ Iterate

# References

## Course:

<http://incompleteideas.net/book/the-book-2nd.html>

<https://lilianweng.github.io/lil-log/2018/02/19/>

<a-long-peek-into-reinforcement-learning.html>

<http://rail.eecs.berkeley.edu/deeprlcourse-fa17/f17docs/>

<https://github.com/kengz/awesome-deep-rl>

## Framework:

<https://spinningup.openai.com/en/latest/>

<https://gym.openai.com/envs/#atari>

<https://github.com/deepmind/bsuite>



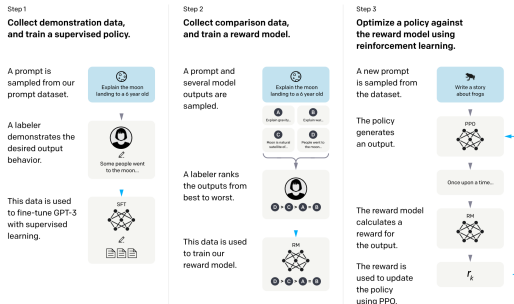


Figure 2: instructgpt-chart-openai