

Introduction to Reinforcement Learning

Model-free

Raphael Cousin - raphael.cousin90@gmail.com
https://github.com/racousin/rl_introduction

Avril 19, 2019

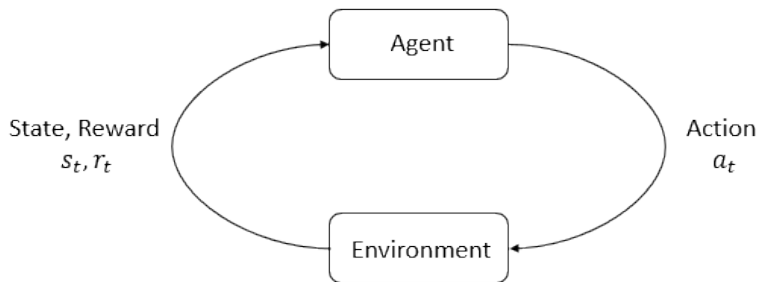
I.1) Introduction

- Supervised learning \rightarrow maps an input to an output based on example input-output pairs.
- Unsupervised learning \rightarrow models the probability density of inputs (using latent variables).
- Reinforcement learning (RL) \rightarrow maps actions to take in an environment to maximize long term reward.

I.2) Introduction

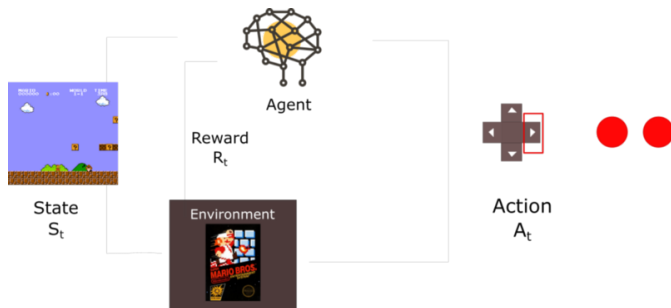
In RL problems:

- Unknown “correct” actions → no explicit supervision for a learning algorithm to try to mimic.
- Provide our algorithms only a reward function → indicates to the learning agent when it is doing well, and when it is doing poorly.



I.3) Introduction

Applications: Neural Architecture Search, Autonomous vehicle, robot legged locomotion, cell-phone network routing, marketing strategy selection, factory control, efficient web-page indexing, video games.



I.4) Markov Decision Process

- S set of all valid states
- A set of all valid actions

- The transition model P

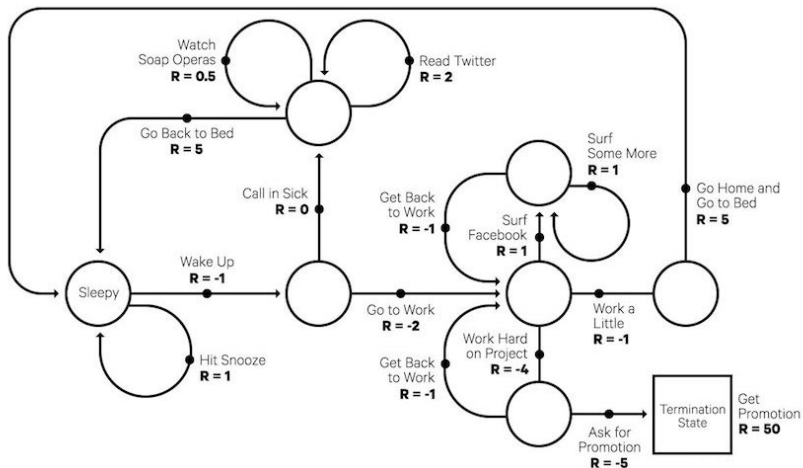
$$P_{ss'}^a = P(s'|s, a) = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} P(s', r | s, a)$$

- The reward function R

$$R(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} P(s', r | s, a)$$

- ρ_0 starting state distribution

I.5) Example: Markov Decision Process



I.6) Policy and trajectory

- Policy (agent brain): rule used to decide what actions to take
 $a_t \sim \pi(\cdot|s_t) \rightarrow \mathbb{P}(a_t|s_t)$
- Trajectories: sequence of states and actions
 $\tau = (s_0, a_0, s_1, a_1, \dots)$
- States follow:
 $s_0 \sim \rho_0(\cdot)$
 $s_{t+1} \sim P(\cdot|s_t, a_t, s_{t-1}, a_{t-1}, \dots, s_0, a_0) = P(\cdot|s_t, a_t)$

I.7) Reward and Return

The future reward, also known as return, is a total sum of rewards going forward.

- Finite-horizon undiscounted return:

$$G_t = \sum_{k=0}^T R_{t+k+1}.$$

- Infinite-horizon discounted return:

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

With $\gamma \in]0, 1[$

I.8) The MDP Objective

- The optimization problem:

$$\pi^* = \arg \max_{\pi} J(\pi)$$

- With the expected return:

$$J(\pi) = E_{\tau \sim \pi}[G(\tau)] = \int_{\tau} P(\tau|\pi)G(\tau)$$

- We know the probability of a T -step trajectory is:

$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$$

I.9) Value Functions

- Value Function, \rightarrow expected return in state s , according to π :

$$V^\pi(s) = E_{\tau \sim \pi}[G_t | s_t = s]$$

- The Optimal Value Function \rightarrow expected return in state s , according to optimal policy:

$$V^*(s) = \max_{\pi} E_{\tau \sim \pi}[G_t | s_t = s] = E_{\pi^*}[G_t | s_t = s]$$

I.10) Action Value Functions

- The Action-Value Function \rightarrow expected return taking action a in state s , and then according π :

$$Q^\pi(s, a) = E_{\tau \sim \pi}[G_t | s_t = s, a_t = a]$$

- The Optimal Action-Value Function \rightarrow expected return taking action a in state s , and then according to optimal policy:

$$Q^*(s, a) = \max_{\pi} E_{\tau \sim \pi}[G_t | s_t = s, a_t = a] = E_{\pi^*}[G_t | s_t = s, a_t = a]$$

I.11) Bellman Equations

- **Idea :** The value of your starting point is the reward you expect to get from being there, plus the value of wherever you land next.

$$V(s) = \mathbb{E}[G_t | S_t = s]$$

$$= \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s]$$

$$= \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$$

$$Q(s, a) = \mathbb{E}[R_{t+1} + \gamma \mathbb{E}_{a \sim \pi} Q(S_{t+1}, a) \mid S_t = s, A_t = a]$$

I.12) Bellman Equations development

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) Q_{\pi}(s, a)$$

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s')$$

$$V_{\pi}(s) = \sum_{a \in \mathcal{A}} \pi(a|s) \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_{\pi}(s') \right)$$

$$Q_{\pi}(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \sum_{a' \in \mathcal{A}} \pi(a'|s') Q_{\pi}(s', a')$$

I.13) Bellman Equations Optimality

Bellman equations for the optimal value functions

$$V_*(s) = \max_{a \in \mathcal{A}} Q_*(s, a)$$

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_*(s')$$

$$V_*(s) = \max_{a \in \mathcal{A}} \left(R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a V_*(s') \right)$$

$$Q_*(s, a) = R(s, a) + \gamma \sum_{s' \in \mathcal{S}} P_{ss'}^a \max_{a' \in \mathcal{A}} Q_*(s', a')$$

I.14) The MDP Solution

Dynamic Programming allows to resolve the MDP optimization problem ($\pi^* = \arg \max_{\pi} E_{\tau \sim \pi}[G(\tau)]$). It is an iterative process:

- Policy evaluation
- Policy improvement

I.15) Policy evaluation

Policy Evaluation: compute the state-value V_π for a given policy π ($\forall s$):

$$\begin{aligned} V_\pi(s) &= \mathbb{E}_\pi[r + \gamma V_\pi(s') | S_t = s] \\ &= \sum_a \pi(a|s) \sum_{s', r} P(s', r|s, a) (r + \gamma V_\pi(s')) \end{aligned}$$

I.16) Policy Improvement

Policy Improvement: generates a better policy $\pi' \geq \pi$ by acting greedily. Compute Q from V ($\forall a, s$):

$$\begin{aligned} Q_{\pi}(s, a) &= \mathbb{E}[R_{t+1} + \gamma V_{\pi}(S_{t+1}) | S_t = s, A_t = a] \\ &= \sum_{s', r} P(s', r | s, a) (r + \gamma V_{\pi}(s')) \end{aligned}$$

Update greedily: $\pi'(s) = \arg \max_{a \in \mathcal{A}} Q_{\pi}(s, a)$ ($\forall s$)

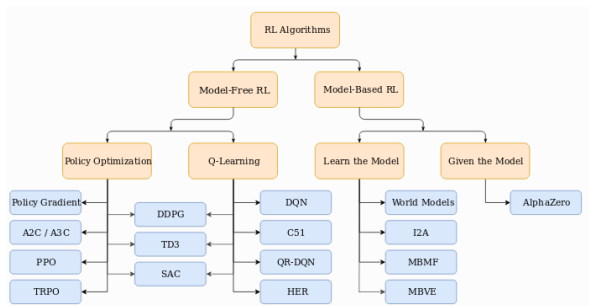
I.17) Dynamic Programming

Policy Iteration: iterative procedure to improve the policy when combining policy evaluation and improvement.

$$\pi_0 \xrightarrow{\text{evaluation}} V_{\pi_0} \xrightarrow{\text{improve}} \pi_1 \xrightarrow{\text{evaluation}} \dots \xrightarrow{\text{improve}} \pi_* \xrightarrow{\text{evaluation}} V_* \quad (1)$$

This policy iteration process works and always converges to the optimality.

I.18) Bestiary of RL Algorithms



To simplify:

- **Model free:** large data + low complexity
- **Model based:** less data + high complexity

II.1) Model-Free

The objective is the same as MDP:

- $J(\pi) = \int_{\tau} P(\tau|\pi)G(\tau) = E_{\tau \sim \pi}[G(\tau)]$
- The optimization problem:
$$\pi^* = \arg \max_{\pi} J(\pi)$$

But we don't know the transition model P . The constraint is therefore to interact intelligently with the environment to obtain the information needed to solve the problem.

- **Q-learning:** learn the action value function Q
 $(\pi'(s) = \arg \max_{a \in \mathcal{A}} Q_{\pi}(s, a))$
- **Policy Optimization:** learn directly the policy π

II.2) Definition

- **On-policy:** Use the deterministic outcomes or samples from the target policy to train the algorithm.
- **Off-policy:** Training on a distribution of transitions or episodes produced by a different behavior policy rather than that produced by the target policy.

II.3) Monte-Carlo

- To evaluate $V(s) = E_{\tau \sim \pi}[G_t | s_t = s]$
- From complete episodes $S_1, A_1, R_2, \dots, S_T$ to compute $G_t = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1}$.
- The empirical value function is : $V(s) = \frac{\sum_{t=1}^T \mathbb{1}[S_t=s] G_t}{\sum_{t=1}^T \mathbb{1}[S_t=s]}$
- As, well, the empirical action-value function is :
- $Q(s, a) = \frac{\sum_{t=1}^T \mathbb{1}[S_t=s, A_t=a] G_t}{\sum_{t=1}^T \mathbb{1}[S_t=s, A_t=a]}$

II.4) Monte-Carlo Algorithm

Initialize Q $Q(s, a) \forall s, a$.

❶ Generate an episode with the policy π (extract from Q ϵ -greedy)

❷ Evaluate Q using the episode:

$$q_{\pi}(s, a) = \frac{\sum_{t=1}^T (\mathbb{1}[S_t=s, A_t=a] \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1})}{\sum_{t=1}^T \mathbb{1}[S_t=s, A_t=a]}$$

❸ Improve the policy greedily: $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$

❹ Iterate

II.5) Temporal difference/bootstrapping

- Remember $V(S_t) = \mathbb{E}[R_{t+1} + \gamma V(S_{t+1}) | S_t = s]$
- So $R_{t+1} + \gamma V(S_{t+1})$ is an unbiased estimate for $V(S_t)$
- As well $R_{t+1} + \gamma Q(S_{t+1}, A_{t+1})$ is an unbiased estimate for $Q(S_t, A_t)$
- $R_{t+1} + \gamma V(S_{t+1})$ is called the TD target.
- *α improvement* :

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

This observation motivates the following algorithm.

II.6) SARSA Algorithm

Initialize Q function $Q(s, a) \forall s, a$

S_t = initial state, act with π to get A_t, R_{t+1}, S_{t+1}

- ❶ Act with π to get $A_{t+1}, R_{t+2}, S_{t+2}$
- ❷ Evaluate Q using the observation step:
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$
- ❸ Improve the policy greedily: $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$
- ❹ $A_t = A_{t+1}, R_{t+1} = R_{t+2}, S_{t+1} = S_{t+2}$. Iterate

II.7) Q-learning

- Remember

$$Q_{\pi^*}(S_t, A_t) = \mathbb{E}[R_{t+1} + \gamma \max_{a'} Q^*(S_{t+1}, a') | S_t = s, A_t = a]$$

- So $R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a')$ is an unbiased estimate for $Q(S_t, A_t)$
- $R_{t+1} + \gamma \max_{a'} Q(S_{t+1}, a')$ is called the Q target.
- α improvement:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$

This observation motivates the following algorithm.

II.8) Q-learning Algorithm

Initialize Q function $Q(s, a) \forall s, a$

S_t = initial state

- ❶ Act with π to get A_t, R_{t+1}, S_{t+1}
- ❷ Evaluate Q using the observation step:
$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_{a \in \mathcal{A}} Q(S_{t+1}, a) - Q(S_t, A_t))$$
- ❸ Improve the policy greedily: $\pi(s) = \arg \max_{a \in \mathcal{A}} Q(s, a)$
- ❹ Iterate

II.9) Policy Optimization

- Parametrization of policy, π_θ .
- We aim to maximize the expected return $J(\pi_\theta) = E_{\tau \sim \pi_\theta}[G(\tau)]$.
- Gradient ascent:
$$\theta_{k+1} = \theta_k + \alpha \nabla_\theta J(\pi_\theta)|_{\theta_k}.$$
- We can proof that:
$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\theta}[\sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) G(\tau)]$$

<https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>

II.10) Monte Carlo Policy gradient estimation

- Set of trajectories $\mathcal{D} = \{\tau_i\}_{i=1,\dots,N}$ from the policy π_θ
- $\hat{\nabla} = \frac{1}{|\mathcal{D}|} \sum_{\tau \in \mathcal{D}} \sum_{t=0}^T \nabla_\theta \log \pi_\theta(a_t|s_t) G_t,$

II.11) Reinforce/VPg algorithm

Initialize policy π_θ

- 1 Generate episodes $\mathcal{D} = \{\tau_i\}_{i=1,\dots,N}$ with the policy π_θ
- 2 Update policy (apply gradient ascent) $\theta \leftarrow \theta + \alpha \hat{\nabla}$
- 3 Iterate

Course:

<http://incompleteideas.net/book/the-book-2nd.html>

<https://lilianweng.github.io/lil-log/2018/02/19/>

<a-long-peek-into-reinforcement-learning.html>

<http://rail.eecs.berkeley.edu/deeprlcourse-fa17/f17docs/>

<https://github.com/kengz/awesome-deep-rl>

Framework:

<https://spinningup.openai.com/en/latest/>

<https://gym.openai.com/envs/#atari>

<https://github.com/deepmind/bsuite>

What we haven't seen

- **AlphaGo** Monte-Carlo-Tree-search
[https://medium.com/applied-data-science/
how-to-build-your-own-alphazero-ai-using-python-and-keras](https://medium.com/applied-data-science/how-to-build-your-own-alphazero-ai-using-python-and-keras)
- **MBRL** <https://arxiv.org/abs/1907.02057>
- **Bandit** <https://arxiv.org/abs/1802.09127>
- **Intrinsic reward** <https://pathak22.github.io/noreward-rl/>